

Facet-Conditioned Trace Audits for Diagnosing Synthetic Persona Failures

Jason Roell

June 25, 2026

Abstract

Synthetic personas increasingly act from long-lived constructed identities rather than single-turn prompts. Their failures often include epistemic components: a system may omit a persona constraint, overweight a conflicting cue, express confidence unsupported by its construction, or provide a plausible but unfaithful justification. This paper introduces the Agentic Epistemology Framework (AEF), a trace schema and audit service for diagnosing such failures. AEF records selected facets, prompt provenance, facet weights, expected behavioral markers, response evidence, calibration checks, and optional counterfactual probes. We extend the framework with *facets*: persistent psychological, socio-ecological, psychographic, and interaction constraints that compose into an agent’s active frame. The reference implementation provides deterministic public harnesses that illustrate which trace fields are needed to localize injected failure classes, plus a preliminary live integration with a Vurvey persona-generation pipeline. In controlled synthetic harnesses, richer traces recover more injected labels than prompt-only or event-only traces; these numbers are best interpreted as trace-sensitivity checks rather than external validity estimates. A small live counterfactual study shows measurable output sensitivity to facet removal, but also reveals the main current limitation: expected-marker metadata is too sparse to diagnose real responses robustly. The contribution is therefore a reproducible audit interface, harness scaffold, and implementation path toward stronger validation of facet-conditioned persona faithfulness.

1 Introduction

Autonomous agents now plan, use tools, communicate, and revise internal state across long-running interactions. Synthetic personas add another layer: they must remain faithful to a constructed identity while responding naturally to novel stimuli. In both cases, important causal factors are often hidden inside prompt text, model activations, unstructured logs, or generated explanations that may not faithfully describe why the model produced a response.

This paper introduces the Agentic Epistemology Framework (AEF), a structured architecture for auditing agent reasoning. AEF represents beliefs as propositions with confidence and justifications, records observer events for epistemic transitions, and models frames as active cognitive perspectives that influence evidence weighting, source trust, and action thresholds. The paper further extends frames with *facets*: typed persona constraints that jointly shape the active frame. Facets can represent psychological dispositions, lived-context constraints, psychographic motivations, and interaction style. In a synthetic human simulation, facets are not independent prompt fragments. They interact nonlinearly, and those interactions should be inspectable.

This first publication of AEF has two goals. First, it defines a vocabulary and implementation pattern for epistemic auditability in agents and synthetic personas. Second, it provides executable harnesses that test whether the resulting traces make known failure causes diagnosable. The framework does not claim that generated persona behavior is valid in the world merely because the system produces plausible text. Instead, it asks whether the system can expose enough structured evidence to debug when a persona fails to follow its own construction.

1.1 Contributions

This paper makes four contributions:

1. A structured epistemic ontology for beliefs, confidence, justifications, frames, facets, perceptions, actions, and observer events.
2. A facet composition model in which persistent persona constraints contribute nonlinearly to an agent’s active frame.
3. A TypeScript reference implementation with deterministic trace-sensitivity harnesses and a preliminary live counterfactual integration.
4. An evaluation protocol that combines structured traces with confidence calibration, black-box perturbation tests, and optional representation-level probes for reasoning faithfulness.

2 Background and Motivation

AEF uses terminology from belief-desire-intention models [Bratman, 1987, Rao and Georgeff, 1991, Wooldridge, 2009], epistemic logic [Fagin et al., 2003], belief revision [Alchourrón et al., 1985, Gärdenfors, 1988], frame semantics [Minsky, 1974, Fillmore, 1976, Lakoff, 1987], argumentation theory [Dung, 1995, Nute, 1994]. The implementation does not claim to satisfy formal AGM revision postulates, argumentation semantics, or a full agent-based society model. These traditions motivate a vocabulary for audit artifacts: what belief was held, what evidence supported it, what frame shaped interpretation, and what changed after new evidence arrived. The paper’s implemented contribution is therefore closer to an engineering audit layer than a formal epistemic calculus.

Recent work shows that natural-language explanations should not be treated as automatically faithful. Chain-of-thought explanations can systematically misrepresent the true reasons for a model’s answer [Turpin et al., 2023], and faithfulness requires intervention-style tests rather than surface plausibility alone [Lanham et al., 2023]. Calibration and self-evaluation research shows that language models can sometimes estimate whether their own answers are correct, but this depends strongly on elicitation format and task setting [Kadavath et al., 2022]. Black-box hallucination detection methods, such as sampling-based consistency checks, provide another way to test whether outputs are stable under repeated generation [Manakul et al., 2023]. Representation engineering and sparse autoencoder work provide tools for probing latent directions or interpretable features when model activations are available [Zou et al., 2023, Cunningham et al., 2023].

AEF implements the structured-trace layer and exposes interfaces where these stronger checks can attach. The current experiments do not run SelfCheck-style sampling, representation probes, or sparse autoencoder analyses. They should be read as trace and runtime-audit studies, not as evidence that all complementary interpretability layers have been integrated.

3 Facet-Conditioned Synthetic Personas

Synthetic human simulation treats a persona as a complex adaptive system rather than a flat list of demographic labels. A persona has internal psychological dispositions, lived-context constraints, motivations, trust heuristics, resource constraints, and linguistic interaction style. In the Vurvey implementation inspected for this work, these variables are represented as a graph of PM facets. A facet is an axis of variation such as income, response style, source trust, or domain expertise. A facet value contains stable identifiers, prompts, taxonomy metadata, and reliability fields. Personas link to selected values through persona fragments. Population molds define facet distributions, mandatory and preferred values, weights, selection modes, and constraints.

This implementation suggests a natural extension to AEF. A **facet** is a typed, persistent constraint with fields:

$$Facet_i = (code_i, value_i, domain_i, prompt_i, weight_i, confidence_i, provenance_i). \quad (1)$$

The active frame is composed from the current stimulus x_t , the persona state S_t , and selected facets $Facet_{1:n}$:

$$F_t = g(S_t, x_t, Facet_{1:n}, I_{1:n}), \quad (2)$$

where $I_{1:n}$ denotes interaction terms between facets. This form is intentionally nonlinear. For example, high extraversion can produce very different behavior when combined with fragile self-confidence and extreme time poverty than when combined with high perceived control and abundant resources.

3.1 Facet Domains

The framework organizes facets into four audit domains:

1. **Psychological facets:** stable traits, anxiety level, self-confidence archetype, and emotional regulation constraints.
2. **Socio-ecological facets:** locale, life stage, socioeconomic status, time poverty, family structure, and other lived-context constraints.
3. **Psychographic facets:** motivations, adoption-curve position, source trust, perceived behavioral control, and decision heuristics.
4. **Interaction facets:** response style, conversational friction, objection tendency, and linguistic modulation.

Each facet should be auditable at three levels: whether it was selected, whether its prompt or structured representation reached the generation path, and whether the final response reflects its expected influence.

4 Epistemic State and Operational Rules

A **belief** is a proposition P held by an agent with confidence $conf \in [0, 1]$ and a justification. A **justification** is evidence, reasoning trace, source, or derivation path supporting a belief or response. A **frame** is the active cognitive lens that modulates interpretation and evidence weighting. A **facet** is a persistent constraint that contributes to frame composition. An **observer event** is a structured record of an epistemic transition.

AEF is guided by five operational rules:

1. **Traceability:** important actions and responses should trace back to beliefs, justifications, frames, and contributing facets.
2. **Explicit uncertainty:** beliefs and evaluator judgments should expose confidence rather than binary assertions.
3. **Frame awareness:** agents should record when frames change evidence weighting, source trust, or interpretation.
4. **Facet faithfulness:** persona responses should expose which facets were expected to constrain the response and which facets were actually expressed.
5. **Intervention readiness:** traces should support counterfactual tests such as removing a facet, swapping a facet value, or perturbing evidence source trust.

5 Confidence, Frame, and Facet Updates

Let $conf_{old}$ denote an agent’s current confidence in proposition P , and let e denote new evidence relevant to P . AEF treats confidence updates as implementation-specific transitions that should be logged, not as a universal belief-revision rule. A simple implementation may interpolate between prior confidence and an evidence-implied target:

$$conf_{new} = (1 - w) \cdot conf_{old} + w \cdot C(e, P), \quad (3)$$

where w is the evidence saliency recorded by the runtime. This equation is a logging contract for a common update pattern, not an AGM-style belief revision model.

For facet-conditioned personas, the current implementation does not fit a parametric interaction model. Instead, it records a facet influence ledger: selected facet values, prompt hashes, configured weights, runtime weights when available, expected markers, observed markers, suppressed facets, and any known interaction terms. The live Vurvey counterfactuals show why this conservative design matters: removing one facet may not remove a behavior when the same constraint is redundantly encoded in the generated biography or adjacent facets. Future work can fit explicit interaction weights, but the implemented contribution is the audit surface needed to discover when such weights would be meaningful.

6 Auditing Faithfulness

AEF distinguishes four kinds of faithfulness evidence.

Trace faithfulness asks whether the system recorded the inputs and transformations it was designed to use. For synthetic personas, this includes selected facets, facet values, prompt fragments, weights, constraints, and response markers.

Behavioral faithfulness asks whether the output reflects the selected persona constraints. This can be measured by entailment checks, contradiction checks, response-style classifiers, and facet-marker coverage.

Counterfactual faithfulness asks whether changing an input that should matter changes the output in the expected direction. Examples include swapping financial outlook from crisis flow to abundance, changing response style from declarer to verbatim repeater, or replacing peer-trust with institutional-trust.

Latent faithfulness asks whether hidden representations align with the claimed frame or facet when activations are available. Representation engineering can test directions corresponding to persona-relevant concepts, and sparse autoencoders can expose interpretable features. These probes are optional because many production systems expose only black-box access, but the framework leaves room for them as stronger evidence.

7 Reference Implementation

The reference implementation is written in TypeScript and is publicly available at <https://github.com/jroell/agent-epistemology-framework>. It implements beliefs, justifications, frames, agents, memory, actions, tools, and observer events. The implementation now includes two trace-sensitivity harnesses: a generic root-cause harness and a facet-frame faithfulness harness.

The Vurvey production codebase informed the facet extension. Its PM facet graph includes global facets and versioned facet values, `headerPrompt` and `detailPrompt` fields used to build persona prompts, `PmPersonaFragment` rows that link personas to selected facet values, mold facet configurations with weights and options, and auto-configuration services that recommend and weight facet distributions for population molds. The current implementation extends that path in two places: `vurvey-python` writes an AEF construction manifest into `ai_personas.facetConfigs` when personas are generated from molds, and `vurvey-api` attaches an AEF audit summary to each simulated survey answer during phase-two generation. These implementation details motivate the audit requirement: a persona response should be diagnosable back to the facet graph, prompt fragments, selected values, construction manifest, composed frame, rendered system prompt, user prompt, and generated response that produced it.

8 Automated Root-Cause Trace-Sensitivity Harness

The generic harness asks whether structured epistemic traces carry the fields needed to identify injected failure labels in scenarios where the ground-truth label is known by construction. It covers four injected causes: frame misweighting, source-trust error, threshold miscalibration, and justification gap. It evaluates four trace conditions: full trace, no-frame trace, no-justification trace, and event-only trace.

We ran the harness with 160 synthetic cases using seed 42. This single run is descriptive; it is not an uncertainty estimate or an external performance claim. The command was:

```
npx ts-node --transpile-only examples/root-cause-benchmark.ts --seed=42 --cases=160
```

The labels in this harness are intentionally injected into trace fields. The table therefore measures information availability under ablations: whether a diagnostic trace schema preserves the evidence a heuristic diagnoser needs. It should not be read as an unbiased estimate of deployed persona-diagnosis accuracy.

Table 1: Generic root-cause trace-sensitivity results. Full traces expose the fields needed to recover all labels in this deterministic synthetic generator. Ablating frame fields, justification/source fields, or epistemic details removes diagnostic capability.

Trace condition	Accuracy	Macro F1	Coverage	Evidence inspected	Trace events
full	1.000	1.000	1.000	8.00	4.00
no_frame	0.750	0.750	0.750	10.50	4.00
no_justification	0.500	0.500	0.500	13.25	4.00
event_only	0.000	0.000	0.000	16.00	4.00




Figure 1: Generic root-cause diagnosis accuracy by trace condition.

9 Facet-Frame Trace-Sensitivity Harness

The second harness targets synthetic persona auditing. Each case constructs a persona from six facets: anxiety level, time poverty, financial outlook, source trust, response style, and adoption-curve position. The harness injects one of five failure types:

1. **Facet omission:** a required facet is absent from the composed frame.
2. **Facet misweighting:** a facet contribution is materially lower or higher than its expected calibration.
3. **Facet conflict:** two facets conflict and the response follows one while ignoring the other.
4. **Persona drift:** the response introduces an identity inconsistent with the facet set.
5. **Response-style violation:** the response ignores the selected interaction-style facet.

The harness evaluates four conditions: full trace, no facet weights, no facet prompts, and event-only trace. We ran 200 cases per seed across 20 deterministic seeds:

```
npx ts-node --transpile-only examples/facet-frame-faithfulness-benchmark.ts \
  --seed=42 --seed-count=20 --cases=200
```

A representative facet-misweighting case illustrates the diagnostic path. The persona was expected to evaluate a premium AI subscription while constrained by crisis-flow financial outlook and a time-poverty facet. The trace showed `time_poverty` with observed weight 0.18 against expected weight 0.78. The response gave a long, careful evaluation despite the time constraint. The diagnostic procedure predicted `facet_misweighting` after inspecting the facet contribution event.

Table 2: Facet-frame faithfulness results. After fixing the omission-control bug, full traces expose all injected labels in this deterministic sensitivity harness across 20 seeds. Removing facet weights harms diagnosis of weighting errors. Removing facet prompts and expected response markers sharply reduces coverage.

Trace condition	Accuracy	95% CI	Macro F1	95% CI	Coverage	95% CI
full	1.000	± 0.000	1.000	± 0.000	1.000	± 0.000
no_facet_markers_weights	0.800	± 0.000	0.772	± 0.003	0.866	± 0.007
no_facet_prompts	0.400	± 0.000	0.400	± 0.000	0.400	± 0.000
event_only	0.000	± 0.000	0.000	± 0.000	0.000	± 0.000




Figure 2: Facet-frame trace-sensitivity accuracy by trace condition.

10 Facet Faithfulness Audit Service Case Study

The harness above motivated a concrete implementation: `FacetFaithfulnessAuditService`. The service takes a persona response, selected facets, rendered prompt provenance, model metadata, self-reported confidence, and optional counterfactual probes. It returns a facet influence ledger, marker coverage, contradictions, counterfactual sensitivity, calibration score, final faithfulness score, primary diagnosis, causal facets, evidence, and recommended debug actions.

The service implements the audit hooks proposed by the framework: selected facet IDs and values, taxonomy versions, prompt hashes, configured and runtime weights, frame composition, suppressed facets, inferred interaction terms, expected markers, observed markers, forbidden marker contradictions, confidence calibration, and counterfactual sensitivity. A Vurvey adapter maps PM facet values, facet codes, `headerPrompt`, `detailPrompt`, confidence scores, taxonomy versions, construction-manifest weights, and rendered prompt provenance into this audit input shape.

We evaluated the audit service on eight hand-constructed Vurvey-style case studies covering faithful responses, prompt omission, facet misweighting, facet conflict, persona drift, response-style violation, and low counterfactual sensitivity. The goal of this experiment is not to prove external predictive validity. It tests whether the implemented service can recover handcrafted fixture labels from increasingly rich evidence conditions.

The experiment exposed an important design lesson. Mean faithfulness score is not the same as diagnostic accuracy. The full audit condition intentionally surfaces more missing markers and calibration penalties than the weaker conditions, so the mean score can remain conservative even while diagnosis improves. This is desirable for an audit system: it should not inflate trust simply because it sees more evidence. The useful output is the diagnosis and debug action, not a single high score.

Table 3: Facet faithfulness audit service results. Response-only and facet-list evidence are insufficient for reliable diagnosis. Adding markers and weights improves diagnosis, and full AEF evidence, including prompt provenance and counterfactual probes, recovers all eight handcrafted fixture labels. These are fixture labels rather than external deployment labels.

Condition	Accuracy	Faithfulness	Marker cov.	Counterfact.	Calibration
response_only	0.125	0.162	0.000	0.500	0.125
facet_list	0.125	0.382	0.000	0.500	0.125
facet_markers_weights	0.750	0.526	0.281	0.500	0.406
full_aef	1.000	0.536	0.281	0.660	0.406




Figure 3: Facet audit diagnosis accuracy by available evidence condition.

In the faithful fixture, the ledger recorded five selected facets, prompt inclusion for each facet, prompt hashes, configured and runtime weights, and interaction terms such as `financial_outlook:crisis_flow` x `adoption_curve_position:innovator`. The service correctly marked the response as faithful while still listing missing secondary markers, such as explicit novelty language. In the prompt omission fixture, the rendered prompt omitted `financial_outlook`; the service diagnosed `facet_prompt_omission` and recommended verifying prompt assembly. In the persona drift fixture, the response introduced an unsupported enterprise CTO identity with unlimited budget; the service diagnosed `persona_drift` and recommended regenerating or constraining the persona biography.

10.1 Calibration Analysis

The hand-constructed audit cases also provide a calibration sanity check. All eight examples have high self-reported confidence, but the measured faithfulness score is much lower. In the high-confidence bin, mean confidence is 0.875 while mean measured faithfulness is 0.536, a gap of 0.339. This is not a deployment calibration result or a full reliability diagram, because the fixture set is small, high-confidence only, and not independently labeled, but it shows that calibration should be analyzed directly rather than left as a table column.

11 Live Vurvey Counterfactual Case Study

The fixture case studies above still use controlled responses. To test the framework against the actual Vurvey persona generation path, we added a Vurvey API counterfactual runner that calls `PersonaAnswerGenerator.generate` for a baseline persona and again for a counterfactual persona variant. The counterfactual variant is produced by applying a facet intervention to the persona’s PM facet values before generation. This case

Table 4: Fixture calibration summary. The high-confidence bin is overconfident relative to measured faithfulness, motivating reliability diagrams and held-out human labels in future evaluations.

Confidence bin	Count	Mean confidence	Mean faithfulness	Gap
0.00-0.33	0	0.000	0.000	0.000
0.34-0.66	0	0.000	0.000	0.000
0.67-1.00	8	0.875	0.536	0.339




Figure 4: Fixture calibration sanity check comparing self-reported confidence with measured faithfulness.

study therefore uses the live survey-simulation prompt path, the selected persona facets, the generated system prompt, and the configured Gemini survey-simulation model rather than simulated counterfactual scores.

We first extracted read-only facet fixtures from staging and production. Staging produced a representative fixture with 25 high-facet-count personas, 15 published molds, and 157 distinct facet codes. Production produced 25 high-facet-count personas and 71 distinct facet codes; the production database currently contains one draft base mold and no mold facet configurations, so production mold distributions are not reported. Common staging facet codes include `adoption_curve`, `agreeableness`, `brand_relationship_orientation`, `caregiving_load`, and `disposable_income`. Common production facet codes include `adoption_curve`, `age`, `agreeableness`, `anxiety_level`, `challenge_mindset`, and several China-specific ecological and commercial context facets.

We then ran the live counterfactual case study on ten staging personas using the stimulus: *Evaluate a premium AI subscription that takes setup time and charges monthly*. Each case selected the first available high-impact facet from `disposable_income`, `financial_outlook`, `time_poverty`, `risk_tolerance`, `adoption_curve`, or `stress_level`, removed that facet, regenerated the response, and compared the baseline and counterfactual outputs.

Table 5: Live Vurvey counterfactual case-study results on ten staging personas. The counterfactual sensitivity is high, indicating that removing selected facets changes generated outputs. Marker coverage remains low, showing that current marker rules are too sparse and literal for real persona responses.

Cases	Faithfulness	Counterfact. sensitivity	Marker cov.	Primary diagnosis
10	0.588	0.852	0.261	facet_selection_gap (9/10)

The qualitative outputs were informative. One staging persona responded: “A monthly cost is already a stretch for anything extra, and if it takes time to set up? Forget about it.” After the `disposable_income` facet was removed, the generated response still referenced money and family constraints, indicating that

similar constraints are redundantly represented in other facets or in the generated biography. Another persona with high disposable income treated the product as plausible if it saved time, while the counterfactual version became more value-conditional. These outcomes are consistent with facets influencing the live generation path, but effects are distributed across multiple identity layers.

The main finding is that real counterfactual sensitivity is measurable in this small case study, but marker coverage and facet redundancy are currently the limiting factors. The audit service correctly avoided overclaiming faithfulness: nine of ten live cases were diagnosed as `facet_selection_gap` because the expected marker catalog did not cover enough of the real facet semantics and because removing one facet did not remove all related behavioral constraints from the biography and adjacent facets. This is a useful failure mode. It shows that the product improvement should be richer facet-level marker metadata, interaction-aware expectations, and persisted construction manifests. The implementation now records such manifests at persona initialization time and uses them at survey-simulation runtime, closing the main traceability gap identified by the live counterfactual experiment.

12 Discussion

The generic harness shows that structured epistemic traces can expose simple injected failures when the trace contains the relevant fields. This is an information-availability result, not a proof that diagnosis is intrinsically hard or easy. The facet harness is more realistic and harder. Full traces recover many failures, but not all. This is useful, but also limiting: structure is necessary for diagnosis but not sufficient for the harder classes of real persona unfaithfulness. It requires tracing facet selection, prompt rendering, frame composition, expected behavioral markers, response content, and sometimes counterfactual or latent evidence. The live case study further shows that facets are not cleanly separable levers: biography text and adjacent facets can redundantly encode the same constraint.

For Vurvey-style synthetic personas, AEF suggests several concrete audit hooks:

1. Record the selected facet IDs, value IDs, codes, taxonomy versions, and prompt hashes used for a response.
2. Record facet weights and interaction terms used to compose the active frame.
3. Record expected response markers derived from selected facets, such as cost objection, peer evidence request, time constraint, or response-style marker.
4. Run counterfactual probes that swap one high-impact facet while holding the rest constant.
5. Track self-rated confidence and compare it with evaluator judgments and response consistency.
6. When model internals are available, add representation-level probes for persona-relevant directions or sparse features.

13 Limitations

The harnesses use controlled synthetic failures with known labels and deterministic heuristic evaluators. They establish that the implementation can carry and expose useful diagnostic fields, not that all synthetic personas are valid or predictive in deployment, and not that the reported synthetic accuracies estimate real-world performance. The generic harness, audit-service case study, and live counterfactual case study are descriptive. The facet harness reports seed variation because it is cheap to run, but it still uses deterministic injected labels. Future work should add noisy traces, hidden labels, independent evaluators, human debugging studies, live product prompt paths, and held-out response data.

Auditability and external validity are separate axes. AEF can show that a persona did or did not follow its own construction; it cannot show that the construction corresponds to real people. External validation requires held-out human response data and validation work such as the broader synthetic-persona validation harness. The framework also depends on the quality of facet definitions, prompt fragments, weights, and constraints. If the facet graph itself encodes weak theory or biased assumptions, AEF will make those assumptions more visible but will not make them correct. Detailed observer logs can expose sensitive persona or respondent data, so production use requires careful retention, access control, and redaction policy.

14 Code and Data Availability

The reference implementation is available at <https://github.com/jroell/agentive-epistemology-framework> under the MIT license. The implementation includes `FacetFaithfulnessAuditService`, the generic root-cause trace-sensitivity harness, the facet-frame trace-sensitivity harness, the facet audit service case study, public harness documentation, and a Vurvey construction-manifest/runtime-audit integration. Generated outputs are written under the repository’s results directory. A Vurvey API branch named `AEF-implementation` contains the product-facing adapter that maps PM facet values and construction-manifest weights into the audit input shape, attaches AEF audit summaries during survey simulation, and includes the live Vurvey counterfactual case-study runner. A Vurvey Python branch named `AEF-implementation` persists AEF construction manifests into `ai_personas.facetConfigs` during mold-based persona generation.

15 Conclusion

The Agentive Epistemology Framework provides an audit architecture for autonomous agents and synthetic personas. By making beliefs, confidence, justifications, frames, facets, and observer events first-class constructs, AEF supports debugging not only what an agent said, but why the system should have expected that behavior from its configured epistemic state. The facet extension connects this audit layer to synthetic human simulation: facets become inspectable contributors to frame composition, and persona faithfulness becomes a more inspectable engineering target rather than a matter of surface plausibility.

References

- Carlos E. Alchourrón, Peter Gärdenfors, and David Makinson. On the logic of theory change: Partial meet contraction and revision functions. *The Journal of Symbolic Logic*, 50(2):510–530, 1985.
- Michael E. Bratman. *Intention, Plans, and Practical Reason*. Harvard University Press, Cambridge, MA, 1987.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. 2023. URL <https://arxiv.org/abs/2309.08600v3>.
- Phan Minh Dung. On the acceptability of arguments and its fundamental role in nonmonotonic reasoning, logic programming and n-person games. *Artificial Intelligence*, 77(2):321–357, 1995.
- Ronald Fagin, Joseph Y. Halpern, Yoram Moses, and Moshe Y. Vardi. *Reasoning About Knowledge*. MIT Press, Cambridge, MA, 2003.
- Charles J. Fillmore. Frame semantics and the nature of language. *Annals of the New York Academy of Sciences*, 280(1):20–32, 1976.
- Peter Gärdenfors. *Knowledge in Flux: Modeling the Dynamics of Epistemic States*. MIT Press, Cambridge, MA, 1988.
- Saurav Kadavath, Tom Conerly, Amanda Aspell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislaw Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared Kaplan. Language models (mostly) know what they know. 2022. URL <https://arxiv.org/abs/2207.05221v4>.
- George Lakoff. *Women, Fire, and Dangerous Things: What Categories Reveal about the Mind*. University of Chicago Press, Chicago, IL, 1987.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamile Lukosiute, Karina Nguyen, Newton Cheng, Nicholas Joseph, Nicholas Schiefer, Oliver Rausch, Robin Larson, Sam McCandlish, Sandipan

- Kundu, Saurav Kadavath, Shannon Yang, Thomas Henighan, Timothy Maxwell, Timothy Telleen-Lawton, Tristan Hume, Zac Hatfield-Dodds, Jared Kaplan, Jan Brauner, Samuel R. Bowman, and Ethan Perez. Measuring faithfulness in chain-of-thought reasoning. 2023. URL <https://arxiv.org/abs/2307.13702v1>.
- Potsawee Manakul, Adian Liusie, and Mark J. F. Gales. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. 2023. URL <https://arxiv.org/abs/2303.08896v3>.
- Marvin Minsky. A framework for representing knowledge. Technical Report AI Memo 306, MIT AI Laboratory, Cambridge, MA, 1974.
- Donald Nute. Defeasible logic. *Handbook of Logic in Artificial Intelligence and Logic Programming*, 3: 353–395, 1994.
- Anand S. Rao and Michael P. Georgeff. Modeling rational agents within a BDI-architecture. In *Proceedings of the 2nd International Conference on Principles of Knowledge Representation and Reasoning*, pages 473–484. Morgan Kaufmann, 1991.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. 2023. URL <https://arxiv.org/abs/2305.04388v2>.
- Michael Wooldridge. *An Introduction to MultiAgent Systems*. John Wiley & Sons, Chichester, UK, 2nd edition, 2009.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xu Wang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to ai transparency. 2023. URL <https://arxiv.org/abs/2310.01405v4>.